

# Deep Learning for Computer Vision (CS231n)

Note: Yunlong Song

Among competing hypotheses, the simplest is the best. — Occam's Razor

## Introduction




Figure 1: LeNet-5.

## Image Classification with Linear Classifiers

### Nearest Neighbor Classifier

- Memorize all training data.
- Predict the label of the most similar training image. Distance:  $L_1$  (Manhattan)/ $L_2$  (Euclidean).

### K-Nearest Neighbor Classifier

- Predict the majority label of the  $k$  most similar training images.
- Hyperparameter:  $k$ .

### Linear Classifier

- $f(x, W) = Wx + b$ .
- $W$ : weights,  $b$ : bias.
- $W \in \mathbb{R}^{D \times C}$ ,  $b \in \mathbb{R}^C$ .
- $D$ : dimension of input,  $C$ : number of classes.

**Multiclass SVM Loss:** Given an example  $(x_i, y_i)$  and using the score  $s = f(x_i, W)$ , the loss is:

$$L_i = \sum_{j \neq y_i} \begin{cases} 0 & \text{if } s_{y_i} \geq s_j + 1 \\ s_j - s_{y_i} + 1 & \text{otherwise} \end{cases}$$

$$= \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

**Softmax Loss:**

$$L_i = -\log P(Y = y_i | X = x_i)$$

$$= -\log \left( \frac{e^{s_{y_i}}}{\sum_j e^{s_j}} \right)$$

## Neural Networks and Backpropagation

## Convolutional Neural Networks

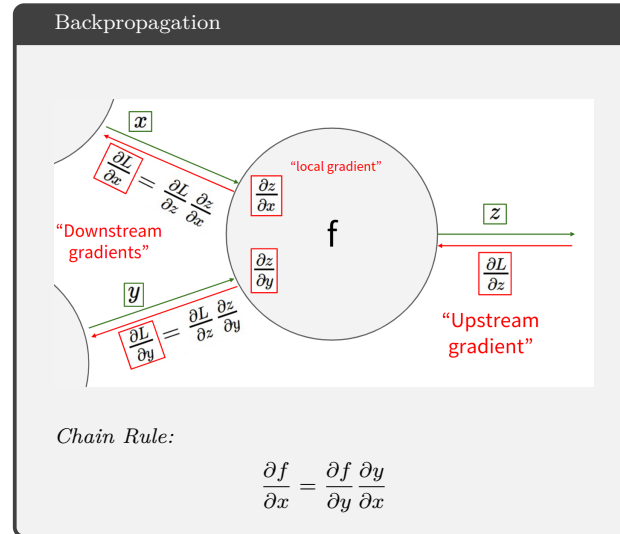
### Neural Networks

- Multiple layers of neurons.
- Convolutional Neural Networks (CNNs).
- Recurrent Neural Networks (RNNs).
- Long Short-Term Memory (LSTM).
- Gated Recurrent Unit (GRU).
- Transformer.

### Activation Functions

- **Sigmoid:**  $\sigma(x) = \frac{1}{1+e^{-x}}$ .
- **Tanh:**  $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$ .
- **ReLU:**  $\text{ReLU}(x) = \max(0, x)$ .
- **Leaky ReLU:**  $\text{LeakyReLU}(x) = \max(0.01x, x)$ .
- **ELU:**  $\text{ELU}(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha(e^x - 1) & \text{if } x \leq 0 \end{cases}$ .

Note: ReLU is the most popular activation function.



A simple example:

$$f(x, y, z) = (x + y)z$$

$$\frac{\partial f}{\partial x} = z, \frac{\partial f}{\partial y} = z, \frac{\partial f}{\partial z} = x + y$$




Figure 2: Kernel Convolution.

### 1D Convolution

$$f(x) * g(x) = \sum_n f(n)g(x - n).$$

### 2D Convolution

$$f(x, y) * g(x, y) = \sum_m \sum_n f(m, n)g(x - m, y - n).$$

### 3D Convolution

$$f(x, y, z) * g(x, y, z) = \sum_m \sum_n \sum_p f(m, n, p)g(x - m, y - n, z - p).$$

### Convolutional Layer

- **Input:**  $W_1 \times H_1 \times D_1$ .
- **Output:**  $W_2 \times H_2 \times D_2$ .
- **Filter:**  $F \times F \times D_1$ .
- **Stride:**  $S$ .
- **Padding:**  $P$ .
- **Output Size:**  $W_2 = \frac{W_1 - F + 2P}{S} + 1$ .
- **Output Size:**  $H_2 = \frac{H_1 - F + 2P}{S} + 1$ .
- **Output Depth:**  $D_2 = \text{Number of filters}$ .

### Pooling Layer

- **Max Pooling:**  $\max(x, y)$ .
- **Average Pooling:**  $\frac{x+y}{2}$ .

# Training Neural Networks

## Activation Function

- **Sigmoid:** Vanishing gradient, not zero-centered.
- **Tanh:** Vanishing gradient, zero-centered.
- **ReLU:** Dying ReLU.
- **Leaky ReLU:** Solves dying ReLU,
- **ELU:** Solves dying ReLU

Note: Use ReLU as default, may try Leaky ReLU or ELU. Don't use sigmoid.

## Data Preprocessing

- **Zero-centered:** Subtract the mean.
- **Normalization:** Divide by the standard deviation.
- **PCA/Whitening:** Decorrelate features.

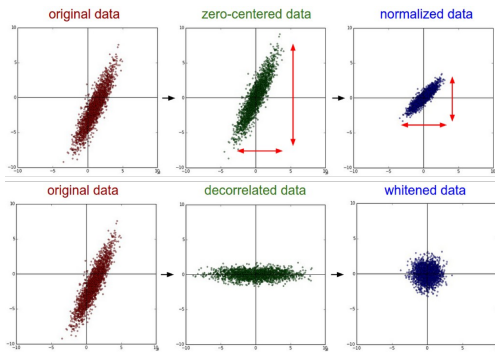


Figure 3: Data Preprocessing.

## Weight Initialization

- **Small random numbers:**  
 $W = 0.01 \times \text{randn}(\text{size}(W))$ .
- **Xavier initialization**
- **He initialization**

**Batch Normalization** Motivation: "you want to normalize the input to each layer so that it has a nice distribution of values, e.g., Gaussian with zero mean and unit variance."

$$\hat{x} = \frac{x^{(k)} - E[x^{(k)}]}{\sqrt{\text{Var}[x^{(k)}] + \epsilon}} \quad (1)$$

## Batch Normalization

**Input:** Values of  $x$  over a mini-batch:  $\{x_1, \dots, x_m\}$ .  
Parameters to learn:  $\gamma, \beta$ .  
**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$ .

$$\mu = \frac{1}{N} \sum_i x_i$$

$$\sigma^2 = \frac{1}{N} \sum_i (x_i - \mu)^2$$

$$\hat{x}_i = \frac{x_i - \mu}{\sqrt{\sigma^2 + \epsilon}}$$

$$y_i = \gamma \hat{x}_i + \beta$$

## Baby Sitting the Learning Process

- **Learning Rate:** Start with a small learning rate.
- **Hyperparameters:** Tune hyperparameters.
- **Monitor Loss:** Plot loss over time.
- **Overfitting:** Regularization, dropout, data augmentation.
- **Visualize:** Visualize weights, activations, gradients.

## Regularization




Figure 4: Regularization.

- $L = \text{data loss} + \text{regularization loss}$ .
- $L = L_i + \lambda R(W)$ .
- $R(W) = \sum_k \sum_l W_{k,l}^2$  (L2 regularization).

## Dropout




Figure 5: Dropout.

## Data Augmentation

- **Random Cropping:** Randomly crop the image.
- **Random Flipping:** Randomly flip the image.
- **Color Jittering:** Randomly change the color.
- **Rotation:** Randomly rotate the image.

## Optimization

- **Random Search:**  
 $W = \text{randn}(\text{size}(W))$ .
- **Gradient Descent:**  
 $W = W - \alpha \nabla_W L$ .
- **Stochastic Gradient Descent (SGD):**  
 $W = W - \alpha \nabla_W L_i$ .
- **Mini-batch Gradient Descent:**  
 $W = W - \alpha \nabla_W L_{\text{batch}}$ .

## SGD with Momentum




Figure 6: SGD with Momentum.

SGD with momentum is a method that helps accelerate SGD in the relevant direction and dampens oscillations. It does this by adding a fraction  $\gamma$  of the update vector of the past time step to the current update vector:

- $v = \gamma v - \alpha \nabla_W L$ .
- $W = W + v$ .

**AdaGrad** AdaGrad is an adaptive learning rate method. The key idea is to adapt the learning rate to the parameters, performing larger updates for infrequent and smaller updates for frequent parameters.

- $G = G + g^2$ .
- $W = W - \alpha \frac{g}{\sqrt{G + \epsilon}}$ .

**RMSProp** RMSProp is an unpublished, adaptive learning rate method proposed by Geoff Hinton. The key idea is to divide the learning rate by an exponentially decaying average of squared gradients.

- $E[g^2]_t = 0.9E[g^2]_{t-1} + 0.1g^2$ .
- $W = W - \alpha \frac{g}{\sqrt{E[g^2] + \epsilon}}$ .

## Adam

Adam looks a bit like RMSProp with momentum. The key idea is to use the first and second moments of the gradients to adaptively adjust the learning rate for each parameter.

- $m = \beta_1 m + (1 - \beta_1)g$ .
- $v = \beta_2 v + (1 - \beta_2)g^2$ .
- $W = W - \alpha \frac{m}{\sqrt{v + \epsilon}}$ .

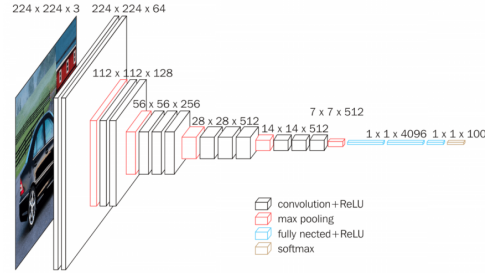



Figure 9: VGGNet.

- 152 layers.
- Skip connections.  
 $y = F(x, \{W_i\}) + x$ .
- Batch normalization.
- SGD with momentum.
- Total parameters: 60 million.


## Recurrent Neural Networks

### Recurrent Neural Networks (RNN)



Hidden state  $h_t$ . Output  $y_t$ .

### Long Short-Term Memory (LSTM)



- Input gate  $i_t$ .
- Forget gate  $f_t$ .
- Output gate  $o_t$ .
- Gate gate (candidate)  $g_t$ .
- Cell state  $c_t$ .
- Hidden state  $h_t$ .

### Gated Recurrent Unit (GRU)

- Update gate:  
 $z_t = \sigma(W_x z x_t + W_h z h_{t-1} + b_z)$ .
- Reset gate:  
 $r_t = \sigma(W_x r x_t + W_h r h_{t-1} + b_r)$ .
- Candidate hidden state:  
 $\tilde{h}_t = \tanh(W_x h x_t + W_h h(r_t \odot h_{t-1}) + b_h)$ .
- Hidden state:  
 $h_t = (1 - z_t) \odot h_{t-1} + z_t \odot \tilde{h}_t$ .

## Deep Learning Software



Figure 7: PyTorch

## CNN Architectures

**AlexNet** [2012, Alex Krizhevsky] The first deep learning model to win the ImageNet Large Scale Visual Recognition Challenge.




Figure 2: An illustration of the architecture of our CNN, explicitly showing the delineation of responsibilities between the two GPUs. One GPU runs the layer-parts at the top of the figure while the other runs the layer-parts at the bottom. The GPUs communicate only at certain layers. The network's input is 150,528-dimensional, and the number of neurons in the network's remaining layers is given by 253,440-186,624-64,896-64,896-43,264-4096-4096-1000.

Figure 8: AlexNet.

- 8 layers: 5 convolutional and 3 fully connected.
- ReLU activation.
- Dropout.
- Data augmentation.
- SGD with momentum.
- Total parameters: 60 million.

**VGGNet** [2014, Karen Simonyan and Andrew Zisserman]

- 19 layers: 16 convolutional and 3 fully connected.
- Small filters:  $3 \times 3$ .
- Max pooling.
- ReLU activation.
- Dropout.
- SGD with momentum.
- Total parameters: 138 million.

**GoogLeNet** [2014, Szegedy et al.]




Figure 10: GoogLeNet.

- 22 layers.
- Inception module.
- Global average pooling.
- SGD with momentum.
- Total parameters: 5 million.

**ResNet** [2015, Kaiming He et al.]





Figure 11: ResNet.

## Sequence-to-Sequence Models

The best YouTube Video on Sequence-to-Sequence Models, including self-attention and transformers. Link: [https://youtu.be/YAgjfMR9R\\_M?si=4Y5vQMbW5fqKONsX](https://youtu.be/YAgjfMR9R_M?si=4Y5vQMbW5fqKONsX). I am too lazy to type the notes; thus, I did many screenshots of the slides. They are great if you know the context.

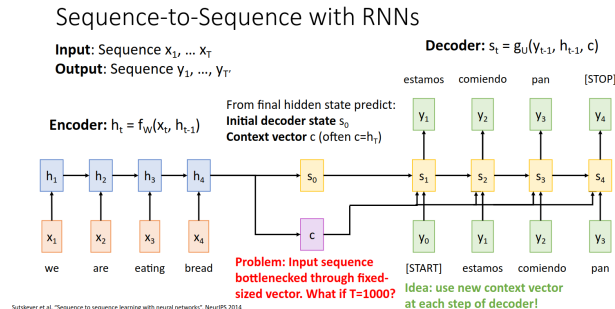
Three ways of processing sequences:

- Recurrent Neural Networks (RNNs).
- 1D Convolutional Neural Networks.
- Self-Attention Networks.

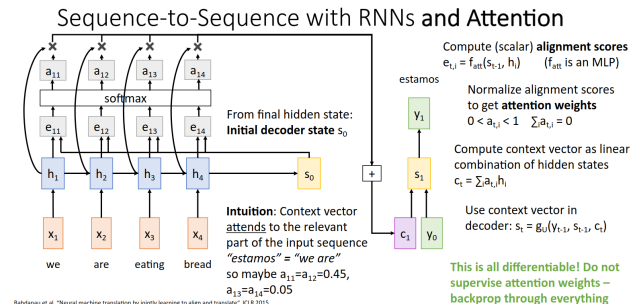


## Attention Mechanism

### Sequence-to-Sequence with RNNs

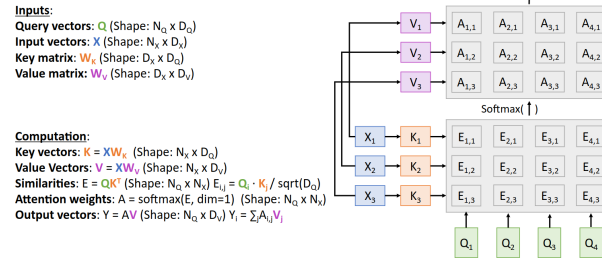


Sutahmet et al. "Sequence to sequence learning with neural networks", NeurIPS 2014



Bahdanau et al. "Neural machine translation by jointly learning to align and translate", ICLR 2015

### Attention Layer



### Self-Attention Layer

One query per input vector




Figure 12: Self-Attention Layer ( The self-attention layer is **Per-mutation Invariant/Equivariant**.).

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

### Self-Attention Layer




Figure 13: Self-Attention Layer with Positional Encoding.

## Transformer

### The Transformer

**Transformer Block:**  
**Input:** Set of vectors  $x$   
**Output:** Set of vectors  $y$

Self-attention is the only interaction between vectors!

Layer norm and MLP work independently per vector

Highly scalable, highly parallelizable

A **Transformer** is a sequence of transformer blocks

Vaswani et al:  
12 blocks,  $D_Q=512$ , 6 heads




Figure 14: Transformer.

## Object Detection

### Loss Function

• **Bounding Box Regression:**  
 $L_{\text{reg}} = \sum_i \sum_j \sum_k \text{smooth}_{L1}(t_{ij}^k - t_{ij}^{k*})$

• **Classification:**  
 $L_{\text{cls}} = - \sum_i \log p_{ij}^*$

• **Total Loss:**  
 $L = L_{\text{reg}} + \lambda L_{\text{cls}}$

- **Sliding Window Detection**
- **Region Proposals:** Selective Search.
- **Region-based CNNs (R-CNN):** Selective Search.
- **Fast R-CNN:** Region of Interest (RoI) Pooling.
- **Faster R-CNN:** Region Proposal Network (RPN).
- **YOLO:** You Only Look Once.
- **SSD:** Single Shot MultiBox Detector.

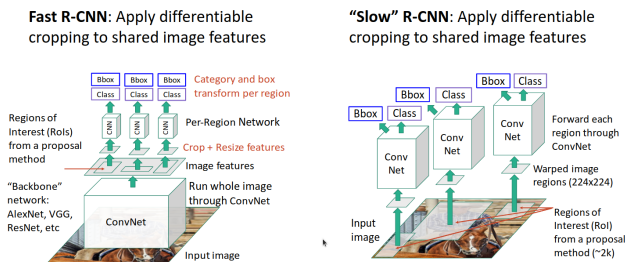


Figure 15: R-CNN

**Intersection over Union (IoU)** How can we compare the predicted bounding box with the ground truth bounding box?

$$\text{IoU} = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$

**Non-Maximum Suppression (NMS)** How can we remove redundant bounding boxes?

- Sort the bounding boxes by confidence.
- Pick the bounding box with the highest confidence.
- Remove all bounding boxes with  $\text{IoU} > \text{threshold}$ .
- Repeat until no more bounding boxes.

**Mean Average Precision (mAP)** How can we evaluate object detection?

- Precision:  $\frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$ .
- Recall:  $\frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$ .




Figure 16: Mean Average Precision.

## Semantic Segmentation

The goal of semantic segmentation is to label each pixel in an image with a corresponding class of what is being represented. The loss function (per pixel cross entropy) is:

$$L = \sum_i \sum_j \text{cross\_entropy}(p_{ij}, p_{ij}^*) = - \sum_i \sum_j p_{ij}^* \log p_{ij}$$

- **Mask R-CNN**: Faster R-CNN + FCN.
- **Segment Anything**

## 3D Vision




Figure 17: 3D Vision Problems




Figure 18: 3D Representation.

- **3D Convolutional Neural Networks.**
- **Volumetric CNNs.**
- **PointNet.**
- **PointNet++.**

### Core Problems

- **Depth Estimation**
- **3D Reconstruction**
- **Pose Estimation**
- **Scene Understanding**

## Video

### Core Problems

- **Action Recognition**
- **Video Classification**
- **Video Segmentation**
- **Video Captioning**

### Video Models

- **Single Frame CNN**
- **Late Fusion**
- **Early Fusion**
- **3D CNN/C3D**
- **Two-Stream Networks**
- **CNN + RNN**
- **Convolutional RNN**
- **Spatio-Temporal self-attention**
- **SlowFast Networks**

## Generative Models

$x$ : data,  $y$ : label.

- **Discriminative Models**:  $P(y|x)$   
Assign a label to an input.  
Feature Learning (supervised learning).
- **Generative Models**:  $P(x)$   
Generate new data.  
Feature Learning (unsupervised learning).  
Detect outliers.

- **Conditional Generative Models**:  $P(x|y)$   
Generate data given a label.  
Assign labels to data, while rejecting outliers.

### Autoregressive Models (Explicit Density Models)

Goal: Write down an explicit function for  $p(x) = f(x, W)$ .  
Given dataset  $\{x_1, \dots, x_N\}$ , maximize the likelihood:

$$\begin{aligned} W^* &= \max_W \prod_i p(x_i) \\ &= \max_W \prod_i f(x_i, W) \\ &= \max_W \sum_i \log f(x_i, W) \\ &= \min_W \sum_i -\log f(x_i, W) \end{aligned}$$

*Example*: PixelRNN and PixelCNN explicitly parameterize density function with a neural network, so we can train to maximize likelihood of training data.

$$p_\theta(x) = \prod_i^n p_\theta(x_i | x_1, \dots, x_{i-1})$$

### Variational Autoencoders (Implicit Density Models)

VAE define an intractable density function that we cannot explicitly compute or optimize. But we will be able to directly optimize a **lower bound** on the likelihood of the data (density).

VAE jointly train an encoder  $q_\phi(z|x)$  and a decoder  $p_\theta(x|z)$  to maximize the variational lower bound on the log-likelihood of the data.

$$\log p_\theta(x) \geq \mathbb{E}_{z \sim q_\phi(z|x)} [\log p_\theta(x|z)] - \text{KL}(q_\phi(z|x) || p(z))$$




Figure 19: Variational Autoencoder.

## Detailed Derivation of VAE

$$\begin{aligned}
 \log p_{\theta}(x) &= \log \frac{p_{\theta}(x|z)p(z)}{p_{\theta}(z|x)} \\
 &= \log \frac{p_{\theta}(x|z)p(z)q_{\phi}(z|x)}{p_{\theta}(z|x)q_{\phi}(z|x)} \\
 &= \log p_{\theta}(x|z) - \log \frac{q_{\phi}(z|x)}{p(z)} + \log \frac{q_{\phi}(z|x)}{p_{\theta}(z|x)} \\
 &= E_z[\log p_{\theta}(x|z)] - E_z\left[\log \frac{q_{\phi}(z|x)}{p(z)}\right] \\
 &\quad + E_z\left[\log \frac{q_{\phi}(z|x)}{p_{\theta}(z|x)}\right] \\
 &= \underbrace{E_z[\log p_{\theta}(x|z)]}_{\text{Reconstruction Loss}} - \underbrace{\text{KL}(q_{\phi}(z|x)||p(z))}_{\text{KL Divergence}} \\
 &\quad + \underbrace{\text{KL}(q_{\phi}(z|x)||p_{\theta}(z|x))}_{\text{cannot compute}} \\
 \log p_{\theta}(x) &\geq \mathbb{E}_{z \sim q_{\phi}(z|x)}[\log p_{\theta}(x|z)] - \text{KL}(q_{\phi}(z|x)||p(z))
 \end{aligned}$$

## Reparameterization Trick

Motivation: During training, we need to compute the gradient of the loss function with respect to the parameters of the encoder. However, the direct sampling process  $z \sim \mathcal{N}(\mu, \sigma^2)$  is not differentiable.

$$\begin{aligned}
 z &= \mu + \sigma \odot \epsilon \\
 \epsilon &\sim \mathcal{N}(0, 1)
 \end{aligned}$$

**Generative Adversarial Networks (Implicit Density Models)** GANs are a framework for estimating generative models via an adversarial process. Different from VAE, GANs do not explicitly model the density function  $p(x)$ , but instead learn a generator function  $G(z)$  that can generate samples.

## GAN

$$\begin{aligned}
 \min_G \max_D L(D, G) &= \mathbb{E}_{x \sim p_{\text{data}}(x)}[\log D(x)] \\
 &\quad + \mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z)))]
 \end{aligned}$$

